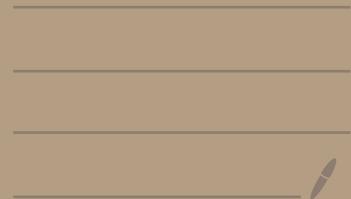


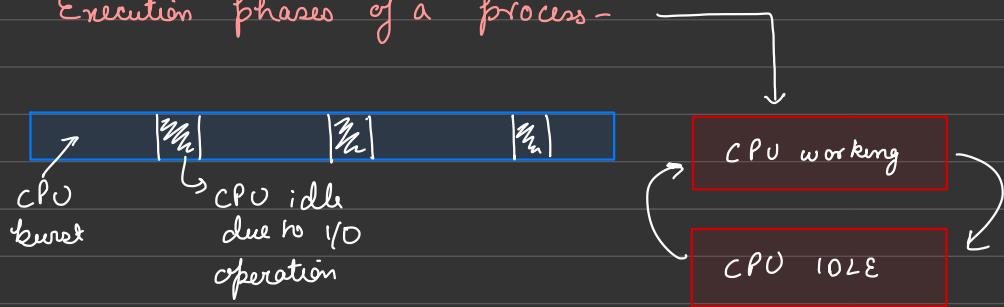
# CPU Scheduling



## # CPU Scheduling

⇒ Which process the scheduler should choose?

Execution phases of a process -



Based on above we can divide process into 2 types -  
(from scheduling perspective)

### ① I/O Bound

has small bursts of CPU activity & then waits for I/O.  
eg - word Processor.

Affects user interaction that's why have higher priorities so  
have to wait for CPU for less time

### ② CPU Bound

Handle any I/O mostly CPU activity (eg 3d rendering)  
have longer CPU burst

Have lower priority.

NOTE → Some process can be both as well ex MS Excel.

## # Scheduling Criteria -

- Max CPU utilization

CPU should not be idle

- Max throughput

Complete as many processes as possible per unit time

- Min turnaround time

For a process, turnaround time, is the time from start to end

- Min response time

Time starts when process goes in the ready queue & ends when CPU picks it & start executing.

CPU should respond immediately.

- Min waiting time

Process should not wait long in the ready queue.  
(sum of all waiting time in ready queue)

- Fairness

Give each process a fair share of CPU.

- Arrival time

When process enters ready queue

## # FCFS (first come first serve)

first job that comes to CPU gets the CPU  
Non preemption

- Process continues till the burst cycle ends

If arrival time is same then choose randomly.

- Advantages → Easy to implement

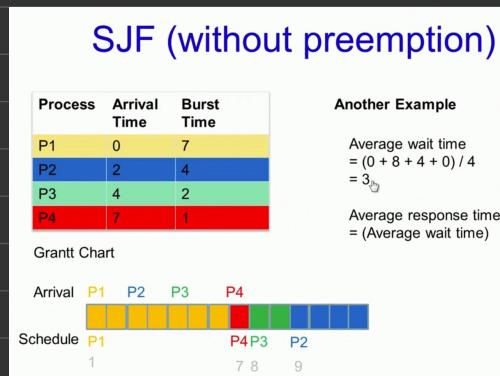
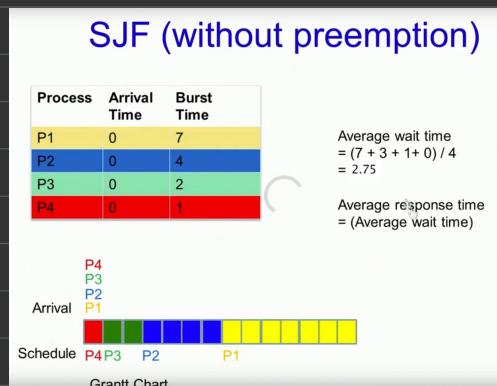
- Disadvantage → Waiting time depends on order

- short process selects waiting for long process to complete

- \* Convoy Effect → All process wait for 1 by process to get off the CPU.

## # Shortest Job First (SJF) { no-preemption }

- Schedules process with the shortest burst time
  - FCFS if same
- No preemption: the process continues to execute until its CPU bursts completes
- SJF, with preemption: the process may get preempted when a new process arrives.

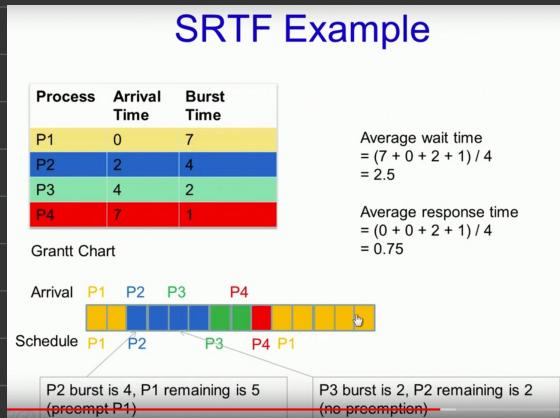


- Advantages →
  - Optimal → Minimum avg. wait time
  - Avg. response time decreases
- Disadvantages →
  - Not practical → difficult to predict burst time
    - Learning to predict future
  - May starve long jobs.

## # Shortest Remaining Time First -- SRTF

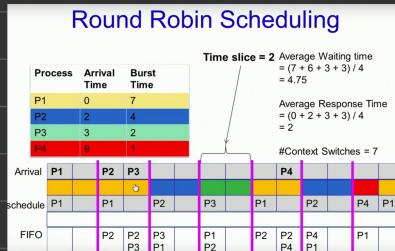
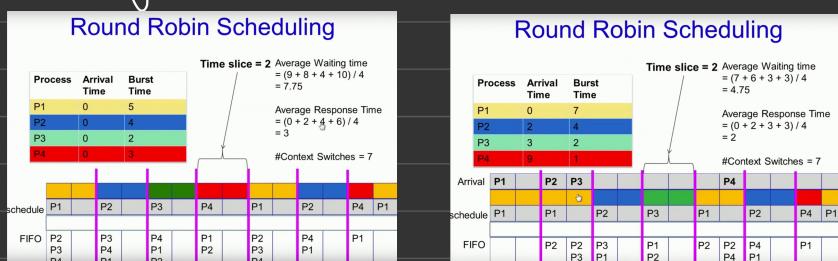
(SJF with preemption)

- If a new process arrives with a shorter burst time than remaining of current process then schedule new process.
- further reduces avg waiting time & avg response time
- Not practical coz you can't predict burst time



## # Round Robin Algorithm

Run process for a time slice then move to back to ready queue.  
 At every timer interrupt preempt current process.



### Example (smaller timeslice)

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	3	2
P4	9	1

Time slice = 1  
 Average Waiting time  
 $= (7 + 7 + 4 + 2) / 4$   
 $= 5$

Average Response Time  
 $= (0 + 1 + 2 + 2) / 4$   
 $= 1.25$

#Context Switches = 11

Arrival	P1	P2	P3	P4											
schedule	P1	P1	P2	P1	P3	P2	P1	P3	P2	P1	P4	P1	P4	P2	P1
FIFO		P2	P3	P3	P2	P1	P3	P2	P1	P4	P4	P2	P1	P1	P3

More context switches but quicker response times

### Example (larger timeslice)

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	3	2
P4	9	1

Time slice = 5  
 Average Waiting time  
 $= (5 + 3 + 9 + 3) / 4$   
 $= 5$

Average Response Time  
 $= (0 + 3 + 6 + 2) / 4$   
 $= 2.75$

#Context Switches = 4

Arrival	P1	P2	P3	P4											
schedule	P1		P2	P3	P1	P4	P2	P1	P3	P1	P4	P3	P2	P1	
FIFO		P2	P3	P3	P1	P3	P1	P3	P1	P1	P4	P3	P1	P4	P3

Lesser context switches but looks more like FCFS (bad response time)

## # Duration of timeslice

### A short quantum

- Good because processes need not wait long before they are scheduled in.
- Bad because context switch overhead increase.

### A long quantum

- Bad because processes no longer seems to execute concurrently they look like FCFS
- May degrade system performance.

Typically kept between 10 ms - 100 ms  
 $\rightarrow \text{avg has } 10 \text{ ms.}$

## # Advantages of Round robin

- faster response time
- low avg wait time when burst time vary.
- fair (Each process gets a fair chance on CPU).

## # Disadvantage of Round Robin

Increased context switching which are overheads.  
 High avg wait times when burst times is equal.

# \* Priority Based Scheduling

Each process is assigned a priority

- A priority is a number in a range (0 - 255)
- A small priority no. means high priority (lower env)

Scheduling policy →

- pick the process from ready queue having highest priority

Advantage -

- Mechanism to provide relative importance to process.

Disadvantage -

- Causes starvation to low priority process.

Starvation		
Process	Arrival Time	Burst Time
P1	0	8
P2	2	4
P3	3	2
P4	9	1

Time slice = 1  
Low priority process may never get a chance to execute.  
P4 is a low priority process

Arrival      P1      P2      P3      P4

schedule      | yellow | blue | grey | yellow | blue | grey | yellow | blue | yellow |

                P1      P2      P3      P4

## # Dealing with starvation

Scheduler adjusts priority of processes to ensure that eventually all of them executes

Some techniques are like -

- Every process is given a base priority  
After every time slot increase the priority of all other processes.
- After process executes, reset priority

# types of priorities →

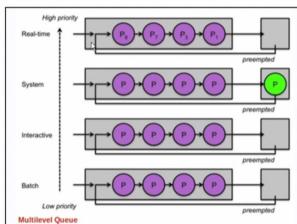
Static Priority → Typically set at start of starvation  
- If not set by user, there is a default priority.

Dynamic Priority → Scheduler can change the process priority during execution in order to achieve scheduling goals  
- e.g. → decrease priority of a process to give another process a chance to execute  
- e.g. → Increase priority for I/O bound process.

# Priority Scheduling with large no. of processes  
Several processes get assigned the same base priority.  
Scheduling begins to behave like round robin

## Multilevel Queues

- Processes assigned to a priority classes
- Each class has its own ready queue
- Scheduler picks the highest priority queue (class) which has at least one ready process
- Selection of a process within the class could have its own policy
  - Typically round robin (but can be changed)
  - High priority classes can implement first come first serve in order to ensure quick response time for critical tasks

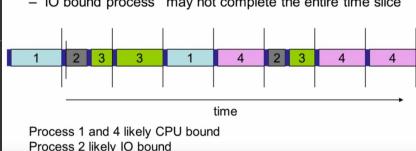


## More on Multilevel Queues

- Scheduler can adjust time slice based on the class picked
  - I/O bound process can be assigned to higher priority classes with longer time slice
  - CPU bound processes can be assigned to lower priority classes with shorter time slices
- Disadvantage :
  - Class of a process must be assigned apriori (not the most efficient way to do things!)

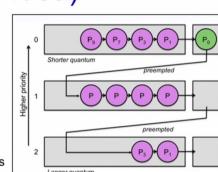
## Multilevel feedback Queues

- Process dynamically moves between priority classes based on its CPU/ IO activity
- Basic observation
  - CPU bound process' likely to complete its entire timeslice
  - IO bound process' may not complete the entire time slice



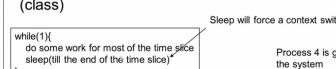
## Multilevel feedback Queues (basic Idea)

- All processes start in the highest priority class
- If it finishes its time slice (likely CPU bound)
  - Move to the next lower priority class
- If it does not finish its time slice (likely IO bound)
  - Keep it on the same priority class
- As with any other priority based scheduling scheme, starvation needs to be dealt with



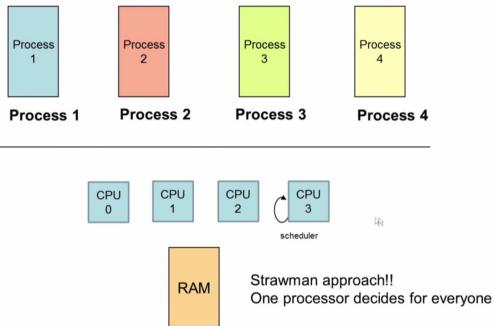
## Gaming the System

- A compute intensive process can trick the scheduler and remain in the high priority queue (class)

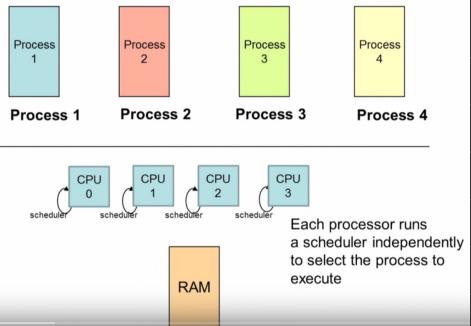


time

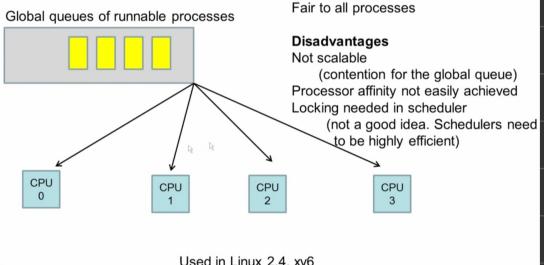
## Multiprocessor Scheduling with a single scheduler



## Multiprocessor Scheduling (Symmetrical Scheduling)



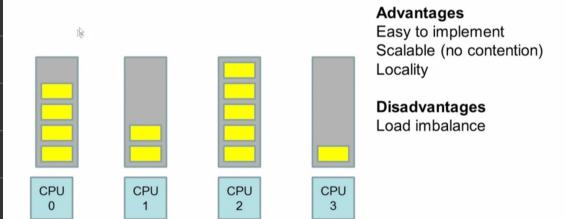
## Symmetrical Scheduling (with global queues)



39

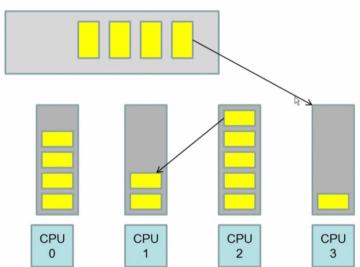
## Symmetrical Scheduling (with per CPU queues)

- Static partition of processes across CPUs



## Hybrid Approach

- Use local and global queues
- Load balancing across queues feasible
- Locality achieved by processor affinity wrt the local queues
- Similar approach followed in Linux 2.6



## Load Balancing

- Two techniques
  - **Push Migration** : A special task periodically monitors load of all processors, and redistributes work when it finds an imbalance
  - **Pull Migration** : Idle processors pull a waiting task from a busy processor
- Process Migration
  - Migration is expensive, it requires all memories to be repopulated